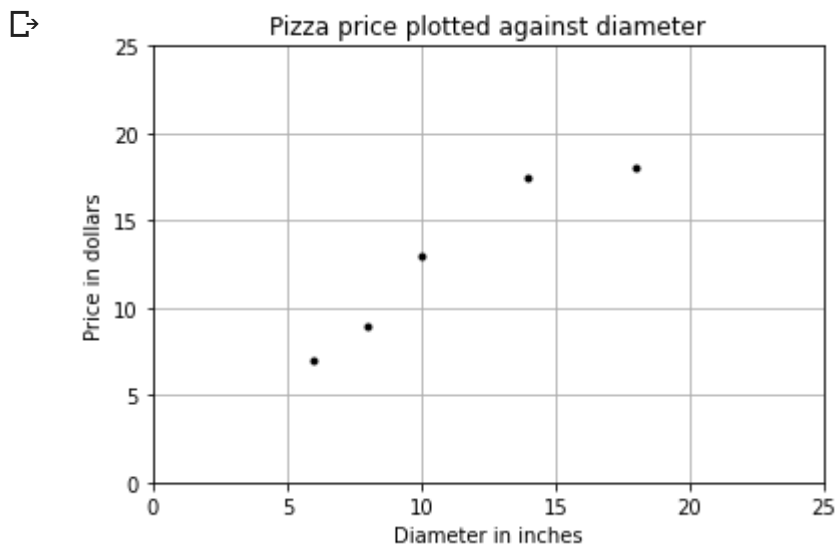


```

1 import numpy as np
2 # "np" and "plt" are common aliases for NumPy and Matplotlib, respectively.
3 import matplotlib.pyplot as plt
4 %matplotlib inline
5 # X represents the features of our training data, the diameters of the pizzas.
6 # A scikit-learn convention is to name the matrix of feature vectors X.
7 # Uppercase letters indicate matrices, and lowercase letters indicate vectors.
8 X = np.array([[6], [8], [10], [14], [18]]).reshape(-1, 1)
9 y = [7, 9, 13, 17.5, 18]
10 # y is a vector representing the prices of the pizzas.
11
12 plt.figure()
13 plt.title('Pizza price plotted against diameter')
14 plt.xlabel('Diameter in inches')
15 plt.ylabel('Price in dollars')
16 plt.plot(X, y, 'k.')
17 plt.axis([0, 25, 0, 25])
18 plt.grid(True)
19 plt.show()

```



```

1 from sklearn.linear_model import LinearRegression
2 model = LinearRegression() # Create an instance of the estimator
3 model.fit(X, y) # Fit the model on the training data
4
5 # Predict the price of a pizza with a diameter that has never been seen before
6 test_pizza = np.array([[12]])
7 predicted_price = model.predict(test_pizza)[0]
8 print('A 12" pizza should cost: $%.2f' % predicted_price)

```

➤ A 12" pizza should cost: \$13.68

```

1 print('Residual sum of squares: %.2f' % np.mean((model.predict(X) - y) ** 2))

```

➤ Residual sum of squares: 1.75

```

1 import numpy as np
2 X = np.array([[6], [8], [10], [14], [18]]).reshape(-1, 1)
3 x_bar = X.mean()
4 print('Mean: ', x_bar)
5
6 # Note that we subtract one from the number of training instances when calculating the sample variance.
7 # This technique is called Bessel's correction. It corrects the bias in the estimation of the population variance
8 # from a sample.
9 variance = ((X - x_bar)**2).sum() / (X.shape[0] - 1)
10 print('Variance: ', variance)
11 # Alternate way
12 print('Variance with Bessel Correction Using ddof: ', np.var(X, ddof=1))

```

➤ Mean: 11.2  
Variance: 23.2  
Variance with Bessel Correction Using ddof: 23.2

```

1 # We previously used a List to represent y.
2 # Here we switch to a NumPy ndarray, which provides a method to calculate the sample mean.
3 y = np.array([7, 9, 13, 17.5, 18])
4
5 y_bar = y.mean()
6 # We transpose X because both operands must be row vectors
7 covariance = np.multiply((X - x_bar).transpose(), y - y_bar).sum() / (X.shape[0] - 1)
8 print('Covariance:', covariance)
9 print('Covariance (Alternate Computation Method): $%.2f' % np.cov(X.transpose(), y)[0][1])

```



```
print( Covariance (Alternate Computation Method): %.2f % np.cov(X.transpose(), y)[0][1])
```

```
➤ Covariance: 22.65
Covariance (Alternate Computation Method): $22.65
```

```
1 import numpy as np
2 from sklearn.linear_model import LinearRegression
3
4 X_train = np.array([6, 8, 10, 14, 18]).reshape(-1, 1)
5 y_train = [7, 9, 13, 17.5, 18]
6
7 X_test = np.array([8, 9, 11, 16, 12]).reshape(-1, 1)
8 y_test = [11, 8.5, 15, 18, 11]
9
10 model = LinearRegression()
11 model.fit(X_train, y_train)
12 r_squared = model.score(X_test, y_test)
13 print('Coefficient of Determination or R2: %.4f' % r_squared )
```

```
➤ Coefficient of Determination or R2: 0.6620
```

Multiple Regression. Three different ways of modeling.

```
1 # In[1]:
2 from numpy.linalg import inv
3 from numpy import dot, transpose
4
5 X = [[1, 6, 2], [1, 8, 1], [1, 10, 0], [1, 14, 2], [1, 18, 0]]
6 y = [[7], [9], [13], [17.5], [18]]
7 print(dot(inv(dot(transpose(X), X)), dot(transpose(X), y)))
```

```
➤ [[1.1875      ]
   [1.01041667]
   [0.39583333]]
```

```
1 from numpy.linalg import lstsq
2
3 X = [[1, 6, 2], [1, 8, 1], [1, 10, 0], [1, 14, 2], [1, 18, 0]]
4 y = [[7], [9], [13], [17.5], [18]]
5 print(lstsq(X, y)[0])
```

```
➤ [[1.1875      ]
   [1.01041667]
   [0.39583333]]
/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:5: FutureWarning: `rcond` parameter will change to the default value of 1e-07 in the future. To use the future default and silence this warning we advise to pass `rcond=None`, to keep using the old, explicitly use `rcond=1e-06`.
```

```
1 # In[1]:
2 from sklearn.linear_model import LinearRegression
3
4 X = [[6, 2], [8, 1], [10, 0], [14, 2], [18, 0]]
5 y = [[7], [9], [13], [17.5], [18]]
6 model = LinearRegression()
7 model.fit(X, y)
8 X_test = [[8, 2], [9, 0], [11, 2], [16, 2], [12, 0]]
9 y_test = [[11], [8.5], [15], [18], [11]]
10 predictions = model.predict(X_test)
11 for i, prediction in enumerate(predictions):
12     print('Predicted: %s, Target: %s' % (prediction, y_test[i]))
13     print('R-squared: %.2f' % model.score(X_test, y_test))
```


```
➤ Predicted: [10.0625], Target: [11]
R-squared: 0.77
Predicted: [10.28125], Target: [8.5]
R-squared: 0.77
Predicted: [13.09375], Target: [15]
R-squared: 0.77
Predicted: [18.14583333], Target: [18]
R-squared: 0.77
Predicted: [13.3125], Target: [11]
R-squared: 0.77
```

Polynomial Regression. We use polynomial regression, a special case of multiple linear regression that models a linear relationship between the response variable and polynomial feature terms. The real-world curvilinear relationship is captured by transforming the features, which are then fit in the same manner as in multiple linear regression.

```

1  # In[1]:
2  import numpy as np
3  import matplotlib.pyplot as plt
4  from sklearn.linear_model import LinearRegression
5  from sklearn.preprocessing import PolynomialFeatures
6
7  X_train = [[6], [8], [10], [14], [18]]
8  y_train = [[7], [9], [13], [17.5], [18]]
9  X_test = [[6], [8], [11], [16]]
10 y_test = [[8], [12], [15], [18]]
11 regressor = LinearRegression()
12 regressor.fit(X_train, y_train)
13 xx = np.linspace(0, 26, 100)
14 yy = regressor.predict(xx.reshape(xx.shape[0], 1))
15 plt.plot(xx, yy)
16 quadratic_featurizer = PolynomialFeatures(degree=2)
17 X_train_quadratic = quadratic_featurizer.fit_transform(X_train)
18 X_test_quadratic = quadratic_featurizer.transform(X_test)
19 regressor_quadratic = LinearRegression()
20 regressor_quadratic.fit(X_train_quadratic, y_train)
21 xx_quadratic = quadratic_featurizer.transform(xx.reshape(xx.shape[0], 1))
22 plt.plot(xx, regressor_quadratic.predict(xx_quadratic), c='r', linestyle='--')
23 plt.title('Pizza price regressed on diameter')
24 plt.xlabel('Diameter in inches')
25 plt.ylabel('Price in dollars')
26 plt.axis([0, 25, 0, 25])
27 plt.grid(True)
28 plt.scatter(X_train, y_train)
29 plt.show()
30 print(X_train)
31 print(X_train_quadratic)
32 print(X_test)
33 print(X_test_quadratic)
34 print('Simple linear regression r-squared: %.4f' % regressor.score(X_test, y_test))
35 print('Quadratic regression r-squared: %.4f' % regressor_quadratic.score(X_test_quadratic, y_test))

```



```
[[6], [8], [10], [14], [18]]
```

```
[[ 1.  6. 36.]
```

```
[ 1.  8. 64.]
```

```
[ 1. 10. 100.]
```

```
[ 1. 14. 196.]
```

```
[ 1. 18. 324.]]
```

```
[[6], [8], [11], [16]]
```

```
[[ 1.  6. 36.]
```

```
[ 1.  8. 64.]
```

```
[ 1. 11. 121.]
```

```
[ 1. 16. 256.]]
```

```
Simple linear regression r-squared:0.8097
```

```
Quadratic regression r-squared: 0.8675
```