

Regression model with sklearn

```
1 # Load all the necessary libraries
2 import pandas as pd
3 import numpy as np
4 from pandas import Series, DataFrame
5 from sklearn.linear_model import LinearRegression
6 from sklearn.model_selection import KFold
7 from sklearn.model_selection import train_test_split
8 from sklearn.metrics import accuracy_score
9
10 # Load the file
11 airport=pd.read_csv('/content/airport.csv')
```



```
1 # Check the number of rows and columns
2 airport.shape
```



```
(19, 13)
```



```
1 # View the dataset
2 print(airport)
```


	year	dep	arr	dep_dem	...	caputil	eff_dep	eff_arr	saer
0	2000	172170	158570	215157	...	46.67	91.93	94.31	93.00
1	2001	147010	133806	176997	...	42.79	93.41	95.48	94.34
2	2002	142779	129089	169072	...	41.13	94.81	95.87	95.29
3	2003	140023	126800	168626	...	41.59	94.26	94.73	94.47
4	2004	161351	160474	203993	...	50.32	93.22	94.93	94.02
5	2005	176738	176139	249869	...	54.73	91.55	93.92	92.63
6	2006	191852	191288	326131	...	57.28	91.98	94.20	92.92
7	2007	222946	222505	407468	...	62.42	90.42	91.82	91.05
8	2008	221203	221458	352677	...	60.47	89.35	92.10	90.67
9	2009	209595	209710	312026	...	57.35	89.91	92.47	91.14
10	2010	200902	200595	277480	...	58.61	92.05	92.89	92.45
11	2011	205231	204685	266431	...	56.99	91.84	92.76	92.30
12	2012	202430	202162	254435	...	54.19	93.28	94.91	94.06
13	2013	203798	203443	264611	...	54.76	93.27	94.07	93.66
14	2014	213113	212793	269706	...	56.86	93.58	93.62	93.60
15	2015	220122	220306	285469	...	59.45	93.68	93.92	93.80
16	2016	225989	226432	296871	...	58.37	93.35	94.09	93.71
17	2017	223126	223879	317699	...	59.47	92.18	93.47	92.81
18	2018	227425	227342	312875	...	59.58	92.29	94.00	93.08


```
[19 rows x 13 columns]
```



```
1 # Create a data frame
2 airport=pd.DataFrame(airport, columns=['caputil', 'dep'])
3 # Define X and Y
4 X=pd.DataFrame(airport['caputil'])
5 Y=pd.DataFrame(airport['dep'])
```



```
1 # Define the model
2 reg = LinearRegression()
```



```
1 # Determine three folds
2 scores=[]
3 kfold = KFold(n_splits=3, shuffle=True, random_state=42)
```



```
1 # Create the iteration for the three folds
2 for i, (train, test) in enumerate(kfold.split(X, Y)):
3     reg.fit(X.iloc[train,:], Y.iloc[train,:])
4     score = reg.score(X.iloc[test,:], Y.iloc[test,:])
5     scores.append(score)
6 print(scores)
```



```
[0.9029142706220207, 0.9048352664783202, 0.5119143545760048]
```



```
1 # Determine the coefficient of determination with sklearn
2 reg = LinearRegression().fit(X, Y)
3 print('Coefficient of Determination:', reg.score(X, Y))
```



```
Coefficient of Determination: 0.8897612752328165
```



```
1 # Determine the coefficient for X
2 print('Coefficients:', reg.coef )
```

```

- print(coefficients_, reg.intercept_)

Coefficients: [[4202.71508638]]

1 # Determine the intercept
2 print('Intercept:', reg.intercept_)

Intercept: [-33354.0929305]

1 # Predict the Y values
2 Y_pred=reg.predict(Y)

1 # Determine the values
2 print('Predicted Values: ',Y_pred)

Predicted Values: [[7.23548102e+08]
 [6.17807791e+08]
 [6.00026103e+08]
 [5.88443420e+08]
 [6.78078928e+08]
 [7.42746105e+08]
 [8.06265941e+08]
 [9.36945164e+08]
 [9.29619831e+08]
 [8.80834714e+08]
 [8.44300512e+08]
 [8.62494066e+08]
 [8.50722261e+08]
 [8.56471575e+08]
 [8.95619866e+08]
 [9.25076696e+08]
 [9.49734026e+08]
 [9.37701652e+08]
 [9.55769124e+08]]


1 # Visualize predicted errors
2 from sklearn.model_selection import cross_val_predict
3 predicted = cross_val_predict(reg, X, Y, cv=10)

1 # Determine the predicted errors
2 print('Predicted Errors: ', predicted)

Predicted Errors: [[160835.55508042]
 [143928.21704009]
 [138393.80845549]
 [140358.20860634]
 [180662.82450213]
 [198806.56295092]
 [209181.70395083]
 [231558.86084681]
 [220561.15973044]
 [207488.73277439]
 [214087.71393934]
 [207154.97715331]
 [193506.03241906]
 [195900.52219425]
 [204795.19251424]
 [215536.69965656]
 [210137.99688364]
 [214613.01298902]
 [216056.74708452]]

1 # Create a graph of the predicted versus cross-validated errors
2 import matplotlib.pyplot as plt
3 %matplotlib inline
4 fig, ax = plt.subplots()
5 ax.scatter(Y, predicted, edgecolors=(0, 0, 0))
6 ax.plot([Y.min(), Y.max()], [Y.min(), Y.max()], 'k--', lw=4)
7 ax.set_xlabel('Measured')
8 ax.set_ylabel('Predicted')
9 plt.title('Prediction by Cross-Validation', fontdict=None, loc='center', pad=None)
10 plt.show()

```



Regression Model with statsmodels

```
|
|
|
```

```
1 # Method 2 Using Statsmodels
2 import statsmodels.api as sm
3 from scipy import stats
```

```
/usr/local/lib/python3.6/dist-packages/statsmodels/tools/_testing.py:19: FutureWarning: pandas.util.testing is deprecated. Use the functions in the public API at
import pandas.util.testing as tm
```

```
1 # Add a constant
2 X2 = sm.add_constant(X)
3 # Define the model
4 est = sm.OLS(Y, X2)
5 # Fit the model
6 est2 = est.fit()
7 print("summary()\n",est2.summary())
8 print("pvalues\n",est2.pvalues)
9 print("tvalues\n",est2.tvalues)
10 print("rsquared\n",est2.rsquared)
11 print("rsquared_adj\n",est2.rsquared_adj)
```

```
summary()
```

OLS Regression Results

```
=====
Dep. Variable:      dep  R-squared:      0.890
Model:              OLS  Adj. R-squared: 0.883
Method:             Least Squares  F-statistic: 137.2
Date:               Sun, 14 Feb 2021  Prob (F-statistic): 1.45e-09
Time:               16:34:20  Log-Likelihood: -201.11
No. Observations:   19  AIC: 406.2
Df Residuals:       17  BIC: 408.1
Df Model:            1
Covariance Type:    nonrobust
=====
```

```
=====
              coef  std err          t  P>|t|  [0.025  0.975]
-----
const  -3.335e+04  1.96e+04   -1.698   0.108  -7.48e+04  8092.445
caputil  4202.7151  358.786   11.714   0.000   3445.742  4959.688
=====
```

```
=====
Omnibus:           2.344  Durbin-Watson:      0.634
Prob(Omnibus):     0.310  Jarque-Bera (JB):      1.898
Skew:              -0.721  Prob(JB):              0.387
Kurtosis:          2.437  Cond. No.              464.
=====
```

Warnings:

```
[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
```

pvalues

```
const  1.077598e-01
```

```
caputil  1.454602e-09
```

```
dtype: float64
```

tvalues

```
const  -1.697874
```

```
caputil  11.713701
```

```
dtype: float64
```

rsquared

```
0.8897612752328165
```

```
rsquared_adj
```

```
0.8832766443641586
```

```
/usr/local/lib/python3.6/dist-packages/scipy/stats/stats.py:1535: UserWarning: kurtosistest only valid for n>=20 ... continuing anyway, n=19
"anyway, n=%i" % int(n))
```

```
1 # Determine predicted values and R-squared with scikitlearn
2 from sklearn.metrics import r2_score
3 predictions = est2.predict(X2)
4
5 print(est2.predict(X2))
6
7 print("r2_score",r2_score(Y,predictions))
```

```
0  162786.620151
```

```
1  146480.085616
```

```
2  139503.578572
```

```
3 141436.827512
4 178126.530216
5 196660.503747
6 207377.427217
7 228979.382761
8 220784.088343
9 207671.617273
10 212967.038282
11 206158.639842
12 194391.037600
13 196786.585199
14 205612.286881
15 216497.318955
16 211958.386661
17 216581.373256
18 217043.671916
dtype: float64
r2_score 0.8897612752328165
```

How to create a dummy variable with sklearn

```
1 import warnings
2 warnings.filterwarnings('ignore')
3 import pandas as pd
4 import numpy as np
5 import seaborn as sns
6 import matplotlib.pyplot as plt
7 df=pd.read_table('http://data.princeton.edu/wws509/datasets/salary.dat', delim_whitespace=True)
8 df.head()
```

	sx	rk	yr	dg	yd	sl
0	male	full	25	doctorate	35	36350
1	male	full	13	doctorate	22	35350
2	male	full	10	doctorate	23	28200
3	female	full	7	doctorate	27	26775
4	male	full	19	masters	30	33696

```
1 # To create a dummy variable
2 dummy=pd.get_dummies(df['sx'])
3 dummy.head()
```

	female	male
0	0	1
1	0	1
2	0	1
3	1	0
4	0	1

```
1 # To merge the dataframes
2 df.merge(dummy, left_index=True, right_index=True)
3 df.head()
```

	sx	rk	yr	dg	yd	sl	female	male
0	male	full	25	doctorate	35	36350	0	1
1	male	full	13	doctorate	22	35350	0	1
2	male	full	10	doctorate	23	28200	0	1
3	female	full	7	doctorate	27	26775	1	0
4	male	full	19	masters	30	33696	0	1

```
1 # To concatenate on column 1
2 df=pd.concat([df, dummy], axis=1)
3 df.head()
```

	sx	rk	yr	dg	yd	sl	female	male
0	male	full	25	doctorate	35	36350	0	1
1	male	full	13	doctorate	22	35350	0	1
2	male	full	10	doctorate	23	28200	0	1

```

1 import pandas as pd
2 df = pd.read_csv('http://statweb.stanford.edu/~tibs/ElemStatLearn/datasets/SAheart.data', index_col=0)# copy data and separate predictors and response
3 X = df.copy()
4 y = X.pop('chd')
5 df.head()

```

	sbp	tobacco	ldl	adiposity	famhist	typea	obesity	alcohol	age	chd
row.names										
1	160	12.00	5.73	23.11	Present	49	25.30	97.20	52	1
2	144	0.01	4.41	28.61	Absent	55	28.87	2.06	63	1
3	118	0.08	3.48	32.28	Present	52	29.14	3.81	46	0
4	170	7.50	6.41	38.03	Present	51	31.99	24.26	58	1
5	134	13.60	3.50	27.78	Present	60	25.99	57.34	49	1

```

1 # compute percentage of chronic heart disease for famhist
2 y.groupby(X.famhist).mean()

```

famhist
Absent 0.237037
Present 0.500000
Name: chd, dtype: float64

Creating a dummy variable with statsmodels

```

1 import statsmodels.formula.api as smf
2 # encode df.famhist as a numeric via pd.Factor
3 est = smf.ols(formula="sbp ~ age+ obesity+adiposity+C(famhist)", data=df).fit()
4 est.summary()

```

OLS Regression Results							
Dep. Variable:	sbp	R-squared:	0.174				
Model:	OLS	Adj. R-squared:	0.167				
Method:	Least Squares	F-statistic:	24.04				
Date:	Thu, 25 Feb 2021	Prob (F-statistic):	4.55e-18				
Time:	11:52:36	Log-Likelihood:	-2006.3				
No. Observations:	462	AIC:	4023.				
Df Residuals:	457	BIC:	4043.				
Df Model:	4						
Covariance Type: nonrobust							
	coef	std err	t	P> t	[0.025	0.975]	
Intercept	104.7562	6.167	16.985	0.000	92.636	116.876	
C(famhist)[T.Present]	-0.6380	1.822	-0.350	0.726	-4.218	2.942	
age	0.4072	0.081	5.028	0.000	0.248	0.566	
obesity	0.2788	0.310	0.900	0.368	-0.330	0.888	
adiposity	0.3596	0.206	1.749	0.081	-0.045	0.764	
Omnibus:	78.787	Durbin-Watson:	1.887				
Prob(Omnibus):	0.000	Jarque-Bera (JB):	140.851				
Skew:	0.989	Prob(JB):	2.60e-31				
Kurtosis:	4.846	Cond. No.	412.				

Warnings:
[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

