

```

1  import numpy as np
2  from scipy import stats
3
4  class BayesLinReg:
5
6      def __init__(self, n_features, alpha, beta):
7          self.n_features = n_features
8          self.alpha = alpha
9          self.beta = beta
10         self.mean = np.zeros(n_features)
11         self.cov_inv = np.identity(n_features) / alpha
12
13     def learn(self, x, y):
14
15         # Update the inverse covariance matrix (Bishop eq. 3.51)
16         cov_inv = self.cov_inv + self.beta * np.outer(x, x)
17
18         # Update the mean vector (Bishop eq. 3.50)
19         cov = np.linalg.inv(cov_inv)
20         mean = cov @ (self.cov_inv @ self.mean + self.beta * y * x)
21
22         self.cov_inv = cov_inv
23         self.mean = mean
24
25         return self
26
27     def predict(self, x):
28
29         # Obtain the predictive mean (Bishop eq. 3.58)
30         y_pred_mean = x @ self.mean
31
32         # Obtain the predictive variance (Bishop eq. 3.59)
33         w_cov = np.linalg.inv(self.cov_inv)
34         y_pred_var = 1 / self.beta + x @ w_cov @ x.T
35
36         return stats.norm(loc=y_pred_mean, scale=y_pred_var ** .5)
37
38     @property
39     def weights_dist(self):
40         cov = np.linalg.inv(self.cov_inv)
41         return stats.multivariate_normal(mean=self.mean, cov=cov)

```

Compute the MAE with Bayes Linear Regression model

```

1  from sklearn import datasets
2  from sklearn import metrics
3
4  X, y = datasets.load_boston(return_X_y=True)
5
6  model = BayesLinReg(n_features=X.shape[1], alpha=.3, beta=1)
7
8  y_pred = np.empty(len(y))
9
10 for i, (xi, yi) in enumerate(zip(X, y)):
11     y_pred[i] = model.predict(xi).mean()
12     model.learn(xi, yi)
13
14 print('Mean Absolute Error:\n', metrics.mean_absolute_error(y, y_pred))

```

```

☞ Mean Absolute Error:
   3.78412506186973

```

Compute MAE with Stochastic Gradient Descent Regressor

```

1  from sklearn import exceptions
2  from sklearn import linear_model
3  from sklearn import preprocessing
4
5  model = linear_model.SGDRegressor(eta0=.15) # here eta0 is the learning rate
6
7  y_pred = np.empty(len(y))
8
9  for i, (xi, yi) in enumerate(zip(preprocessing.scale(X), y)):
10     try:
11         y_pred[i] = model.predict([xi])[0]

```

```

65 ax.set_title(f'Posterior target distribution #{i + 1}')
66 # Plot the old points and the new points
67 ax.scatter([xi[1] for xi in xs[:-1]], ys[:-1])
68 ax.scatter(xs[-1][1], ys[-1], marker='*')
69 # Plot the predictive mean along with the predictive interval
70 ax.plot(w, [p.mean() for p in posteriors], linestyle='--')
71 cis = [p.interval(.95) for p in posteriors]
72 ax.fill_between(
73     x=w,
74     y1=[ci[0] for ci in cis],
75     y2=[ci[1] for ci in cis],
76     alpha=.1
77 )
78 # Plot the true target distribution
79 ax.plot(w, [np.dot(weights, [1, xi]) for xi in w], color='red')


```



```

12     except exceptions.NotFittedError:
13         y_pred[i] = 0.
14     model.partial_fit([xi], [yi])
15
16 print('Mean Absolute Error:\n', metrics.mean_absolute_error(y, y_pred))

```

 Mean Absolute Error:

4.171800265184048

In a Bayesian linear regression, the weights follow a distribution that quantifies their uncertainty. In the case where there are two features – and therefore two weights in a linear regression – this distribution can be represented with a contour plot. As for the predictive distribution, which quantifies the uncertainty of the model regarding the spread of possible feature values, we can visualize it with a shaded area, as is sometimes done in control charts.

```

1  from mpl_toolkits.axes_grid1 import ImageGrid
2  import matplotlib.pyplot as plt
3  %matplotlib inline
4  np.random.seed(42)
5
6  # Pick some true parameters that the model has to find
7  weights = np.array([-0.3, 0.5])
8
9  def sample(n):
10     for _ in range(n):
11         x = np.array([1, np.random.uniform(-1, 1)])
12         y = np.dot(weights, x) + np.random.normal(0, 0.2)
13         yield x, y
14
15  model = BayesLinReg(n_features=2, alpha=2, beta=25)
16
17  # The following 3 variables are just here for plotting purposes
18  N = 100
19  w = np.linspace(-1, 1, 100)
20  W = np.dstack(np.meshgrid(w, w))
21
22  n_samples = 5
23  fig = plt.figure(figsize=(7 * n_samples, 21))
24  grid = ImageGrid(
25      fig, 111, # similar to subplot(111)
26      nrows_ncols=(n_samples, 3), # creates a n_samplesx3 grid of axes
27      axes_pad=0.5 # pad between axes in inch.
28  )
29
30  # We'll store the features and targets for plotting purposes
31  xs = []
32  ys = []
33
34  def prettify_ax(ax):
35      ax.set_xlim(-1, 1)
36      ax.set_ylim(-1, 1)
37      ax.set_xlabel('$w_1$')
38      ax.set_ylabel('$w_2$')
39      return ax
40
41  for i, (xi, yi) in enumerate(sample(n_samples)):
42
43      pred_dist = model.predict(xi)
44
45      # Prior weight distribution
46      ax = prettify_ax(grid[3 * i])
47      ax.set_title(f'Prior weight distribution #{i + 1}')
48      ax.contourf(w, w, model.weights_dist.pdf(W), N, cmap='viridis')
49      ax.scatter(*weights, color='red') # true weights the model has to find
50
51      # Update model
52      model.learn(xi, yi)
53
54      # Posterior weight distribution
55      ax = prettify_ax(grid[3 * i + 1])
56      ax.set_title(f'Posterior weight distribution #{i + 1}')
57      ax.contourf(w, w, model.weights_dist.pdf(W), N, cmap='viridis')
58      ax.scatter(*weights, color='red') # true weights the model has to find
59
60      # Posterior target distribution
61      xs.append(xi)
62      ys.append(yi)
63      posteriors = [model.predict(np.array([1, wi])) for wi in w]
64      ax = prettify_ax(grid[3 * i + 2])

```

