

```

1 import numpy as np # linear algebra
2 import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)

1 # Matplotlib and seaborn for visualization
2 import matplotlib.pyplot as plt
3 %matplotlib inline
4
5 import seaborn as sns
6
7 # Linear Regression to verify implementation
8 from sklearn.linear_model import LinearRegression
9
10 # Scipy for statistics
11 import scipy
12
13 # PyMC3 for Bayesian Inference
14 import pymc3 as pm

```

```

1 exercise = pd.read_csv('/content/exercise.csv')
2 calories = pd.read_csv('/content/calories.csv')
3 df = pd.merge(exercise, calories, on = 'User_ID')
4 df = df[df['Calories'] < 300]
5 df = df.reset_index()
6 df['Intercept'] = 1
7 df.head()

```

```

↳

```

	index	User_ID	Gender	Age	Height	Weight	Duration	Heart_Rate	Body_Temp	Calories	Intercept
	0	14733363	male	68	190.0	94.0	29.0	105.0	40.8	231	1
	1	14861698	female	20	166.0	60.0	14.0	94.0	40.3	66	1
	2	11179863	male	69	179.0	79.0	5.0	88.0	38.7	26	1
	3	16180408	female	34	179.0	71.0	13.0	100.0	40.5	71	1
	4	17771927	female	27	154.0	58.0	10.0	81.0	39.8	35	1

```

1 # Create the features and response
2 X = df.loc[:, ['Intercept', 'Duration']]
3 y = df.loc[:, 'Calories']

1 from sklearn.linear_model import LinearRegression
2 reg = LinearRegression().fit(X, y)
3 print('Coefficient of determination or R2:\n',round(reg.score(X, y),3))

```

```

↳ Coefficient of determination or R2:
0.913

```

```

1 print('Coefficients: \n', reg.coef_)

```

```

↳ Coefficients:
[0.          7.16978335]

```

```

1 print('Intercept:\n', reg.intercept_)

```

```

↳ Intercept:
-21.8281025260508

```

Build the model with 500 observations and then all observations.

```

1 with pm.Model() as linear_model_500:
2     # Intercept
3     intercept = pm.Normal('Intercept', mu = 0, sd = 10)
4
5     # Slope
6     slope = pm.Normal('slope', mu = 0, sd = 10)
7
8     # Standard deviation
9     sigma = pm.HalfNormal('sigma', sd = 10)
10
11     # Estimate of mean
12     mean = intercept + slope * X.loc[0:499, 'Duration']
13
14     # Observed values

```

```

14     # Observed values
15     Y_obs = pm.Normal('Y_obs', mu = mean, sd = sigma, observed = y.values[0:500])
16
17     # Sampler
18     step = pm.NUTS()
19
20     # Posterior distribution
21     linear_trace_500 = pm.sample(1000, step)

```

[]> INFO:pymc3:Sequential sampling (2 chains in 1 job)
 INFO:pymc3:NUTS: [sigma, slope, Intercept]
 100%|██████████| 1500/1500 [00:02<00:00, 642.71it/s]
 100%|██████████| 1500/1500 [00:02<00:00, 611.80it/s]
 WARNING:pymc3:The acceptance probability does not match the target. It is 0.8803147835931578, but should be close to
 WARNING:pymc3:The acceptance probability does not match the target. It is 0.9061663130361706, but should be close to

```

1  with pm.Model() as linear_model:
2      # Intercept
3      intercept = pm.Normal('Intercept', mu = 0, sd = 10)
4
5      # Slope
6      slope = pm.Normal('slope', mu = 0, sd = 10)
7
8      # Standard deviation
9      sigma = pm.HalfNormal('sigma', sd = 10)
10
11     # Estimate of mean
12     mean = intercept + slope * X.loc[:, 'Duration']
13
14     # Observed values
15     Y_obs = pm.Normal('Y_obs', mu = mean, sd = sigma, observed = y.values)
16
17     # Sampler
18     step = pm.NUTS()
19
20     # Posterior distribution
21     linear_trace = pm.sample(1000, step)

```

[]> INFO:pymc3:Sequential sampling (2 chains in 1 job)
 INFO:pymc3:NUTS: [sigma, slope, Intercept]
 100%|██████████| 1500/1500 [00:12<00:00, 117.01it/s]
 100%|██████████| 1500/1500 [00:09<00:00, 159.42it/s]
 WARNING:pymc3:The acceptance probability does not match the target. It is 0.9895795767051766, but should be close to
 WARNING:pymc3:The acceptance probability does not match the target. It is 0.9790896525672083, but should be close to

Compute posterior prediction for 15.5 minutes

```

1  bayes_prediction = linear_trace['Intercept'] + linear_trace['slope'] * 15.5

```

```

1  print('Prediction:\n:', bayes_prediction)

```

[]> Prediction:
 : [89.28474968 89.30393246 89.33247202 ... 89.35222443 89.23276362
 89.21184103]