



A problem with linear regression is that estimated coefficients can become large, making the model sensitive to inputs and possibly unstable. This is particularly true for problems with few observations (samples) or more input predictors (p) than variables than samples (n).

One approach to address the stability of regression models is to change the loss function to include additional costs for a model that has large coefficients. Linear regression models that use these modified loss functions during training are referred to collectively as penalized linear regression.

Least Angle Regression, LAR or LARS for short, is an alternative approach to solving the optimization problem of fitting the penalized model. Technically, LARS is a forward stepwise version of feature selection for regression that can be adapted for the Lasso model.

Unlike the Lasso, it does not require a hyperparameter that controls the weighting of the penalty in the loss function. Instead, the weighting is discovered automatically by LARS.

A problem with linear regression is that estimated coefficients of the model can become large, making the model sensitive to inputs and possibly unstable. This is particularly true for problems with few observations (samples) or more input predictors (p) than variables than samples (n).

One approach to address the stability of regression models is to change the loss function to include additional costs for a model that has large coefficients. Linear regression models that use these modified loss functions during training are referred to collectively as penalized linear regression.

Least Angle Regression, LAR or LARS for short, is an alternative approach to solving the optimization problem of fitting the penalized model. Technically, LARS is a forward stepwise version of feature selection for regression that can be adapted for the Lasso model.

Unlike the Lasso, it does not require a hyperparameter that controls the weighting of the penalty in the loss function. Instead, the weighting is discovered automatically by LARS.

```
1 # load and summarize the housing dataset
2 from pandas import read_csv
3 from matplotlib import pyplot
4 # load dataset
5 url = 'https://raw.githubusercontent.com/jbrownlee/Datasets/master/housing.csv'
6 dataframe = read_csv(url, header=None)
7 # summarize shape
8 print(dataframe.shape)
9 # summarize first few lines
10 print(dataframe.head())
```

```
(506, 14)
```

```
0      1      2      3      4      5      ...      8      9      10      11      12      13
```

0	0.00632	18.0	2.31	0	0.538	6.575	...	1	296.0	15.3	396.90	4.98	24.0
1	0.02731	0.0	7.07	0	0.469	6.421	...	2	242.0	17.8	396.90	9.14	21.6
2	0.02729	0.0	7.07	0	0.469	7.185	...	2	242.0	17.8	392.83	4.03	34.7
3	0.03237	0.0	2.18	0	0.458	6.998	...	3	222.0	18.7	394.63	2.94	33.4
4	0.06905	0.0	2.18	0	0.458	7.147	...	3	222.0	18.7	396.90	5.33	36.2

[5 rows x 14 columns]

```

1 # define model
2 from sklearn.linear_model import Lars
3 model = Lars()

1 # evaluate an lars regression model on the dataset
2 from numpy import mean
3 from numpy import std
4 from numpy import absolute
5 from pandas import read_csv
6 from sklearn.model_selection import cross_val_score
7 from sklearn.model_selection import RepeatedKFold
8 from sklearn.linear_model import Lars
9 # load the dataset
10 url = 'https://raw.githubusercontent.com/jbrownlee/Datasets/master/housing.csv'
11 dataframe = read_csv(url, header=None)
12 data = dataframe.values
13 X, y = data[:, :-1], data[:, -1]
14 # define model
15 model = Lars()
16 # define model evaluation method
17 cv = RepeatedKFold(n_splits=10, n_repeats=3, random_state=1)
18 # evaluate model
19 scores = cross_val_score(model, X, y, scoring='neg_mean_absolute_error', cv=cv, n_jobs=-1)
20 # force scores to be positive
21 scores = absolute(scores)
22 print('Mean MAE: %.3f (%.3f)' % (mean(scores), std(scores)))

```

Mean MAE: 3.432 (0.552)

```

1 # make a prediction with a lars regression model on the dataset
2 from pandas import read_csv
3 from sklearn.linear_model import Lars
4 # load the dataset
5 url = 'https://raw.githubusercontent.com/jbrownlee/Datasets/master/housing.csv'
6 dataframe = read_csv(url, header=None)
7 data = dataframe.values
8 X, y = data[:, :-1], data[:, -1]
9 # define model
10 model = Lars()
11 # fit model
12 model.fit(X, y)
13 # define new data

```

```
14 row = [0.00032, 18.00, 2.310, 0, 0.5380, 0.5750, 65.20, 4.0900, 1, 290.0, 15.30, 390.90, 4.98]
15 # make a prediction
16 yhat = model.predict([row])
17 # summarize prediction
18 print('Predicted: %.3f' % yhat)
```

Predicted: 29.904

```
1 # use automatically configured the lars regression algorithm
2 from numpy import arange
3 from pandas import read_csv
4 from sklearn.linear_model import LarsCV
5 from sklearn.model_selection import RepeatedKFold
6 # load the dataset
7 url = 'https://raw.githubusercontent.com/jbrownlee/Datasets/master/housing.csv'
8 dataframe = read_csv(url, header=None)
9 data = dataframe.values
10 X, y = data[:, :-1], data[:, -1]
11 # define model evaluation method
12 cv = RepeatedKFold(n_splits=10, n_repeats=3, random_state=1)
13 # define model
14 model = LarsCV(cv=cv, n_jobs=-1)
15 # fit model
16 model.fit(X, y)
17 # summarize chosen configuration
18 print('alpha: %f' % model.alpha_)
```

alpha: 0.001623