

One-Class Support Vector Machine | Example 1

```
1 # import libraries
2 import pandas as pd
3 from sklearn.svm import OneClassSVM
4 import matplotlib.pyplot as plt
5 from numpy import where

1 # import data
2 data = pd.read_csv("https://raw.githubusercontent.com/uiuc-cse/data-fa14/gh-pages/data/iris.csv")
3 # input data
4 df = data[["sepal_length", "sepal_width"]]

1 # model specification
2 # nu = 0.03 means that the algorithm will designate 3% data as outliers.
3 model = OneClassSVM(kernel = 'rbf', gamma = 0.001, nu = 0.03).fit(df)

1 # prediction
2 y_pred = model.predict(df)
3 y_pred

array([ 1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1, -1,  1,  1,  1,
        1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,
        1,  1,  1,  1,  1,  1,  1, -1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,
        1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,
        1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,
        1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,
        1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1, -1, -1,
        1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1, -1,  1,  1,  1,  1,
        1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1,  1])

1 # filter outlier index
2 outlier_index = where(y_pred == -1)
3 # filter outlier values
4 outlier_values = df.iloc[outlier_index]
5 outlier_values
```


	sepal_length	sepal_width
13	4.3	3.0
41	4.5	2.3
117	7.7	3.8
118	7.7	2.6
131	7.9	3.8

```

1 # visualize outputs
2 plt.scatter(data["sepal_length"], df["sepal_width"])
3 plt.scatter(outlier_values["sepal_length"], outlier_values["sepal_width"], c = "r")
4 plt.title("Visualizing Anomalies")

```

Text(0.5, 1.0, 'Visualizing Anomalies')



One-Class Support Vector Machine | Example 2

```

1 # one-class svm for imbalanced binary classification
2 from sklearn.datasets import make_classification
3 from sklearn.model_selection import train_test_split
4 from sklearn.metrics import f1_score, classification_report, confusion_matrix, plot_confusion_matrix
5 from sklearn.svm import OneClassSVM
6 # generate dataset
7 X, y = make_classification(n_samples=10000, n_features=2, n_redundant=0,
8     n_clusters_per_class=1, weights=[0.999], flip_y=0, random_state=4)
9 # split into train/test sets
10 trainX, testX, trainy, testy = train_test_split(X, y, test_size=0.5, random_state=2, stratify=y)
11 # define outlier detection model
12 model = OneClassSVM(gamma='scale', nu=0.01)
13 # fit on majority class
14 trainX = trainX[trainy==0]
15 model.fit(trainX)
16 # detect outliers in the test set
17 yhat = model.predict(testX)
18 # mark inliers 1, outliers -1
19 testy[testy == 1] = -1
20 testy[testy == 0] = 1
21 # calculate F1 score
22 score = f1_score(testy, yhat, pos_label=-1)
23 print('F1 Score: %.3f' % score)
24 # Classification report
25 target_names = ['class 0', 'class 1']
26 print(classification_report(testy, yhat, target_names=target_names))
27 print('Classification Report: ', classification_report)

```

```
27 print(Classification Report: , classification_report)
```

F1 Score: 0.123

	precision	recall	f1-score	support
class 0	0.07	0.80	0.12	5
class 1	1.00	0.99	0.99	4995
accuracy			0.99	5000
macro avg	0.53	0.89	0.56	5000
weighted avg	1.00	0.99	0.99	5000

Classification Report: <function classification_report at 0x7f0976cb48c0>

```
1 # Provide the confusion matrix
2 cm = confusion_matrix(testy, yhat)
3 cm[:, :-1, :-1]
```

```
array([[4939, 56],
       [ 1, 4]])
```

Local Outlier Factor Method. A simple approach to identifying outliers is to locate those examples that are far from the other examples in the feature space.

This can work well for feature spaces with low dimensionality (few features), although it can become less reliable as the number of features is increased, referred to as the curse of dimensionality.

The local outlier factor, or LOF for short, is a technique that attempts to harness the idea of nearest neighbors for outlier detection. Each example is assigned a scoring of how isolated or how likely it is to be outliers based on the size of its local neighborhood. Those examples with the largest score are more likely to be outliers.

```
1 # local outlier factor for imbalanced classification
2 from numpy import vstack
3 from sklearn.datasets import make_classification
4 from sklearn.model_selection import train_test_split
5 from sklearn.metrics import f1_score
6 from sklearn.neighbors import LocalOutlierFactor
7
8 # make a prediction with a lof model
9 def lof_predict(model, trainX, testX):
10     # create one large dataset
11     composite = vstack((trainX, testX))
12     # make prediction on composite dataset
13     yhat = model.fit_predict(composite)
14     # return just the predictions on the test set
15     return yhat[len(trainX):]
16
17 # generate dataset
```

```

18 X, y = make_classification(n_samples=10000, n_features=2, n_redundant=0,
19   n_clusters_per_class=1, weights=[0.999], flip_y=0, random_state=4)
20 # split into train/test sets
21 trainX, testX, trainy, testy = train_test_split(X, y, test_size=0.5, random_state=2, stratify=y)
22 # define outlier detection model
23 model = LocalOutlierFactor(contamination=0.01)
24 # get examples for just the majority class
25 trainX = trainX[trainy==0]
26 # detect outliers in the test set
27 yhat = lof_predict(model, trainX, testX)
28 # mark inliers 1, outliers -1
29 testy[testy == 1] = -1
30 testy[testy == 0] = 1
31 # calculate score
32 score = f1_score(testy, yhat, pos_label=-1)
33 print('F1 Score: %.3f' % score)
34 # Classification report
35 target_names = ['class 0', 'class 1']
36 print(classification_report(testy, yhat, target_names=target_names))
37 print('Classification Report: ', classification_report)

```

F1 Score: 0.138

	precision	recall	f1-score	support
class 0	0.08	0.80	0.14	5
class 1	1.00	0.99	0.99	4995
accuracy			0.99	5000
macro avg	0.54	0.90	0.57	5000
weighted avg	1.00	0.99	0.99	5000

Classification Report: <function classification_report at 0x7f0976cb48c0>

