

```
1 import pandas as pd
2 import numpy as np
3 from pandas import DataFrame
4 from sklearn.linear_model import LogisticRegression, LinearRegression
5 from sklearn.svm import SVC
6 from sklearn.tree import DecisionTreeClassifier, DecisionTreeRegressor
7 from sklearn.neighbors import KNeighborsClassifier
8 from sklearn.ensemble import VotingClassifier, BaggingRegressor, RandomForestClassifier, ExtraTreesClassifier
9 from sklearn.feature_selection import SelectPercentile
10 from sklearn.preprocessing import RobustScaler, KBinsDiscretizer, StandardScaler, OneHotEncoder
11 from sklearn.compose import ColumnTransformer
12 from sklearn.model_selection import cross_val_score, KFold
13 from sklearn.datasets import load_breast_cancer
```

Voting Classifiers

```
1 #dataset
2 cancer=load_breast_cancer()
3 cancer_data =cancer.data
4 cancer_target =cancer.target
5
6 #classifiers
7 lr = LogisticRegression()
8 dt = DecisionTreeClassifier()
9 svm= SVC(probability=True)
10 knn= KNeighborsClassifier()
11 hard_voting = VotingClassifier(estimators=[('lr', lr), ('dt', dt), ('svc', svm),
12                                     ('knn',knn)],voting='hard')
13 soft_voting = VotingClassifier(estimators=[('lr', lr), ('dt', dt), ('svc', svm),
14                                     ('knn',knn)],voting='soft')
15
16 #data preprocessing
17 anova=SelectPercentile(percentile=30)
18 scale=RobustScaler()
19 ct=ColumnTransformer([('scale',scale,list(range(30))))])
20
21 cancer_data=ct.fit_transform(cancer_data)
22 cancer_data=anova.fit_transform(cancer_data,cancer_target)
23
24 #comparison
25 estimators = [lr,dt,svm,knn,hard_voting,soft_voting]
26 for x in estimators:
27     result=cross_val_score(x,cancer_data,cancer_target,cv=5)
28     final_result=np.mean(result)
29     print(x,final_result)
30
31 #KNeighborsClassifier
32 knn=KNeighborsClassifier(max_iter=-1, probability=True, random_state=None, shrinking=True, tol=0.001,
33                         verbose=False) 0.9437820214252446
34 KNeighborsClassifier(algorithm='auto', leaf_size=30, metric='minkowski',
35                     metric_params=None, n_jobs=None, n_neighbors=5, p=2,
36                     weights='uniform') 0.9455519329296692
37
38 #VotingClassifier
39 voting=VotingClassifier(estimators=[('lr',
40                                     LogisticRegression(C=1.0, class_weight=None,
41                                                         dual=False, fit_intercept=True,
42                                                         intercept_scaling=1,
43                                                         l1_ratio=None, max_iter=100,
44                                                         multi_class='auto',
45                                                         n_jobs=None, penalty='l2',
46                                                         random_state=None,
47                                                         solver='lbfgs', tol=0.0001,
48                                                         verbose=0, warm_start=False)),
49                                   ('dt',
50                                     DecisionTreeClassifier(ccp_alpha=0.0,
51                                                             class_weight=None,
52                                                             criterion='gini',...
53                                                             decision_function_shape='ovr', degree=3,
54                                                             gamma='scale', kernel='rbf', max_iter=-1,
55                                                             probability=True, random_state=None,
56                                                             shrinking=True, tol=0.001, verbose=False)),
57                                   ('knn',
58                                     KNeighborsClassifier(algorithm='auto',
59                                                         leaf_size=30,
60                                                         metric='minkowski',
61                                                         metric_params=None,
62                                                         n_jobs=None, n_neighbors=5,
63                                                         p=2, weights='uniform'))]),
64                             flatten_transform=True, n_jobs=None, voting='hard',
65                             weights=None) 0.9525694767893185
66
67 #VotingClassifier
68 voting=VotingClassifier(estimators=[('lr',
69                                     LogisticRegression(C=1.0, class_weight=None,
70                                                         dual=False, fit_intercept=True,
71                                                         intercept_scaling=1,
72                                                         l1_ratio=None, max_iter=100,
73                                                         multi_class='auto',
74                                                         n_jobs=None, penalty='l2',
75                                                         random_state=None,
76                                                         solver='lbfgs', tol=0.0001,
77                                                         verbose=0, warm_start=False)),
78                                   ('dt',
79                                     DecisionTreeClassifier(ccp_alpha=0.0,
80                                                             class_weight=None,
81                                                             criterion='gini',...
82                                                             decision_function_shape='ovr', degree=3,
83                                                             gamma='scale', kernel='rbf', max_iter=-1,
84                                                             probability=True, random_state=None,
85                                                             shrinking=True, tol=0.001, verbose=False)),
86                                   ('knn',
```

```

KNeighborsClassifier(algorithm='auto',
                    leaf_size=30,
                    metric='minkowski',
                    metric_params=None,
                    n_jobs=None, n_neighbors=5,
                    p=2, weights='uniform'))],
flatten_transform=True, n_jobs=None, voting='soft',
weights=None) 0.947290793355069

```

Bagging and Pasting

```

1  kf = KFold(n_splits=3)
2  bagging = BaggingClassifier(SVC(), n_estimators=500,
3                             max_samples=0.8, bootstrap=True, n_jobs=-1)
4
5  bagging_result = cross_val_score(bagging,cancer_data,cancer_target,cv=kf,n_jobs=-1)
6  print("bagging results",bagging_result)
7  print("average of bagging:",np.mean(bagging_result))
8  print("*****")
9
10 pasting = BaggingClassifier(SVC(), n_estimators=500,
11                             max_samples=0.8, bootstrap=False, n_jobs=-1)
12
13 pasting_result = cross_val_score(pasting,cancer_data,cancer_target,cv=kf,n_jobs=-1)
14 print("pasting results",pasting_result)
15 print("average of pasting:",np.mean(pasting_result))
16 print("*****")

```

```

bagging results [0.89473684 0.97368421 0.96296296]
average of bagging: 0.9437946718648473
*****

pasting results [0.89473684 0.97368421 0.96296296]
average of pasting: 0.9437946718648473
*****

```

With Regressors

```

1  #dataset
2  carset=pd.read_csv('/content/CarPrice_Assignment.csv')
3  carset=carset.drop(['car_ID','symboling','CarName','doornumber','carbody','enginelocation'],axis=1)
4  car_target=carset['price']
5  car_data=carset.iloc[:,0:19]
6  #data preprocessing
7  kf = KFold(n_splits=4)
8  bins = KBinsDiscretizer(n_bins=5, encode='onehot-dense', strategy='uniform')
9  numeric_cols=car_data.select_dtypes(include=np.number).columns
10 print("numeric_cols",numeric_cols)
11 numeric_cols=numeric_cols.delete([1,2,3])
12 categorical_cols=car_data.select_dtypes(exclude=np.number).columns
13 print("categorical_cols",categorical_cols)
14
15 ct=ColumnTransformer([('scaling',StandardScaler(),numeric_cols),
16                        ('binning',bins,['carheight','carwidth','carlength']),
17                        ('categorical',OneHotEncoder(sparse=False,handle_unknown='ignore'),categorical_cols)
18                        ])
19
20 final_data=ct.fit_transform(car_data)
21
22 #bagging_regressor
23 bagging_regressor=BaggingRegressor(DecisionTreeRegressor(), n_estimators=250,
24                                    max_samples=0.8,max_features=0.4, bootstrap=True, n_jobs=-1)
25
26
27 bagging_regressor_result = cross_val_score(bagging_regressor,final_data,car_target,cv=kf,n_jobs=-1)
28 print("bagging results",bagging_regressor_result)
29 print("average of bagging:",np.mean(bagging_regressor_result))
30 print("*****")

```

```

numeric_cols Index(['wheelbase', 'carlength', 'carwidth', 'carheight', 'curbweight',
                   'enginesize', 'boreratio', 'stroke', 'compressionratio', 'horsepower',
                   'peakrpm', 'citympg', 'highwaympg'],
                   dtype='object')
categorical_cols Index(['fueltype', 'aspiration', 'drivewheel', 'enginetype', 'cylindernumber',
                       'fuelsystem'],
                       dtype='object')
bagging results [0.8848483  0.82759289 0.75087733 0.75584312]
average of bagging: 0.8047904103494231
*****

```

Random Forest Classifier

```

1  kf = KFold(n_splits=3)
2
3  random_forest = RandomForestClassifier(n_estimators=250, max_depth=7, n_jobs=-1)
4  rf_results = cross_val_score(random_forest,cancer_data,cancer_target,cv=kf,n_jobs=-1)
5  print("random forest results:",rf_results)
6  print("average of rf:",np.mean(rf_results))
7  print("*****")
8
9  bagging_rf = BaggingClassifier(random_forest,n_estimators=250,max_samples=0.8,bootstrap=True,n_jobs=-1)
10 bagging_rf_result = cross_val_score(bagging_rf,cancer_data,cancer_target,cv=kf,n_jobs=-1)
11 print(" bagging random forest results:",bagging_rf_result)

```

```
12 print("average of bagging rf:",np.mean(bagging_rf_result))
13 print("*****")

random forest results: [0.92631579 0.97368421 0.95767196]
average of rf: 0.9525573192239859
*****

bagging random forest results: [0.91052632 0.97368421 0.96296296]
average of bagging rf: 0.9490578297595841
*****
```

Extra Trees

```
1 kf = KFold(n_splits=3)
2 extra_tree =ExtraTreesClassifier(n_estimators=250,max_depth=7, bootstrap=True,n_jobs=-1)
3 extra_tree_result = cross_val_score(extra_tree,cancer_data,cancer_target,cv=kf,n_jobs=-1)
4 print(" extra tree results:",extra_tree_result)
5 print("average of extra tree:",np.mean(extra_tree_result))

extra tree results: [0.90526316 0.98421053 0.96296296]
average of extra tree: 0.9508122157244964
```