



```
1 from sklearn.datasets import load_wine
2 import pandas as pd
3 import numpy as np
4 np.set_printoptions(precision=4)
5 from matplotlib import pyplot as plt
6 import seaborn as sns
7 sns.set()
8 from sklearn.preprocessing import LabelEncoder
9 from sklearn.tree import DecisionTreeClassifier
10 from sklearn.model_selection import train_test_split
11 from sklearn.metrics import confusion_matrix
```

```
1 wine = load_wine()
2 X = pd.DataFrame(wine.data, columns=wine.feature_names)
3 y = pd.Categorical.from_codes(wine.target, wine.target_names)
```

```
1 X.shape
```


 (178, 13)

```
1 X.head()
```



	alcohol	malic_acid	ash	alcalinity_of_ash	magnesium	total_phenols	flavanoids	nonflavanoid_phenols	pr
0	14.23	1.71	2.43	15.6	127.0	2.80	3.06	0.28	
1	13.20	1.78	2.14	11.2	100.0	2.65	2.76	0.26	
2	13.16	2.36	2.67	18.6	101.0	2.80	3.24	0.30	
3	14.37	1.95	2.50	16.8	113.0	3.85	3.49	0.24	
4	13.24	2.59	2.87	21.0	118.0	2.80	2.69	0.39	


```
1 wine.target_names
```

 array(['class_0', 'class_1', 'class_2'], dtype='<U7')

```
1 df = X.join(pd.Series(y, name='class'))
```

Linear Discriminant Analysis can be broken up into the following steps: Compute the within class and between class scatter matrices
Compute the eigenvectors and corresponding eigenvalues for the scatter matrices
Sort the eigenvalues and select the top k
Create a new matrix containing eigenvectors that map to the k eigenvalues
Obtain the new features (i.e. LDA components) by taking the dot product of the data and the matrix from step 4. For every class, we create a vector with the means of each feature.

```
1 class_feature_means = pd.DataFrame(columns=wine.target_names)
2 for c, rows in df.groupby('class'):
3     class_feature_means[c] = rows.mean()
4 class_feature_means
```



	class_0	class_1	class_2
alcohol	13.744746	12.278732	13.153750
malic_acid	2.010678	1.932676	3.333750
ash	2.455593	2.244789	2.437083
alcalinity_of_ash	17.037288	20.238028	21.416667
magnesium	106.338983	94.549296	99.312500
total_phenols	2.840169	2.258873	1.678750
flavanoids	2.982373	2.080845	0.781458
nonflavanoid_phenols	0.290000	0.363662	0.447500
proanthocyanins	1.899322	1.630282	1.153542
color_intensity	5.528305	3.086620	7.396250
hue	1.062034	1.056282	0.682708
od280/od315_of_diluted_wines	3.157797	2.785352	1.683542
proline	1115.711864	519.507042	629.895833

Then, we plug the mean vectors (mi) into the equation from before in order to obtain the within class scatter matrix.


```
1  within_class_scatter_matrix = np.zeros((13,13))
2  for c, rows in df.groupby('class'):
3      rows = rows.drop(['class'], axis=1)
4      s = np.zeros((13,13))
5  for index, row in rows.iterrows():
6      x, mc = row.values.reshape(13,1), class_feature_means[c].values.reshape(13,1)
7      s += (x - mc).dot((x - mc).T)
8      within_class_scatter_matrix += s
```

```
1  feature_means = df.mean()
2  between_class_scatter_matrix = np.zeros((13,13))
3  for c in class_feature_means:
4      n = len(df.loc[df['class'] == c].index)
5
6      mc, m = class_feature_means[c].values.reshape(13,1), feature_means.values.reshape(13,1)
7
8      between_class_scatter_matrix += n * (mc - m).dot((mc - m).T)
```

```
1  eigen_values, eigen_vectors = np.linalg.eig(np.linalg.inv(within_class_scatter_matrix).dot(between_class_scatter_matrix))
```


The eigenvectors with the highest eigenvalues carry the most information about the distribution of the data. Thus, we sort the eigenvalues from highest to lowest and select the first k eigenvectors. In order to ensure that the eigenvalue maps to the same eigenvector after sorting, we place them in a temporary array.

```
1  pairs = [(np.abs(eigen_values[i]), eigen_vectors[:,i]) for i in range(len(eigen_values))]
2  pairs = sorted(pairs, key=lambda x: x[0], reverse=True)
3  for pair in pairs:
4      print(pair[0])
```

 9.884546449232964
2.9033610617160606
6.285916968291436e-16
6.285916968291436e-16
5.979482586809227e-16
5.345289989557e-16
5.345289989557e-16
2.5624197998855253e-16
2.5624197998855253e-16
2.4778227778444637e-16
5.049704088349899e-17
1.0023889228649853e-17
0.0

It is difficult to determine how much of the variance is explained by each component. Thus, we express it as a percentage.

```
1  eigen_value_sums = sum(eigen_values)
2  print('Explained Variance')
3  for i, pair in enumerate(pairs):
4      print('Eigenvector {}: {}'.format(i, (pair[0]/eigen_value_sums).real))
```

 Explained Variance
Eigenvector 0: 0.772960426932225
Eigenvector 1: 0.22703957306777506
Eigenvector 2: 4.915516446228146e-17
Eigenvector 3: 4.915516446228146e-17
Eigenvector 4: 4.6758882027333914e-17
Eigenvector 5: 4.179956716906465e-17
Eigenvector 6: 4.179956716906465e-17
Eigenvector 7: 2.0037834944392412e-17
Eigenvector 8: 2.0037834944392412e-17
Eigenvector 9: 1.9376295736600758e-17
Eigenvector 10: 3.948811863102963e-18
Eigenvector 11: 7.838568757294643e-19
Eigenvector 12: 0.0

First, we create a matrix W with the first two eigenvectors.

```
1  w_matrix = np.hstack((pairs[0][1].reshape(13,1), pairs[1][1].reshape(13,1))).real
2  w_matrix
```



```
array([[ -4.1562e-02,  2.0879e-01],
       [-1.6684e-02,  3.6500e-03],
       [-4.4776e-03,  4.5911e-01],
       [ 7.5359e-03, -2.4950e-02],
       [ 6.4335e-03,  3.1071e-03],
       [ 9.7213e-02,  1.9083e-01],
```

Then, we save the dot product of X and W into a new matrix Y. where X is a n×d matrix with n samples and d dimensions, and Y is a n×k matrix with n samples and k (k<n) dimensions. In other words, Y is composed of the LDA components, or said yet another way, the new feature space.

```
1  X_lda = np.array(X.dot(w_matrix))
2  X_lda
```




[- - - - - , - - - - -],
[-2.6703, 2.4536],
[-4.0763, 2.5482],
[-2.1484, 2.3438],
[-2.5665, 2.342],
[-2.8082, 1.6951],
[-2.5682, 2.3082],
[-2.7733, 2.5182],
[-2.2102, 2.4526],
[-2.3226, 2.1558],
[-2.0354, 2.7213],
[-1.0681, 3.5209],
[-1.1498, 3.3567],
[-1.08 , 3.2958],
[-1.0542, 3.4031],
[-1.033 , 3.3645],
[-1.1067, 3.4793],
[-1.0068, 3.5505],
[-1.0075, 3.2089],
[-1.0538, 3.3798],
[-0.985 , 3.3684],
[-1.0918, 3.3882],
[-1.1791, 3.4601],
[-1.0039, 3.4614],
[-1.2402, 3.3053],
[-1.0357, 3.3117],
[-1.0564, 3.4578],
[-0.8066, 3.4534],
[-1.0643, 3.5227],
[-1.0027, 3.5641],
[-1.1572, 3.4807],
[-1.1166, 3.5132],
[-0.9805, 3.3335],
[-0.8516, 3.3555],
[-1.076 , 3.4411],
[-0.8448, 3.2103],
[-1.073 , 3.4821],
[-1.039 , 3.3952],
[-1.0692, 3.5956],
[-1.2492, 3.3961],
[-1.1306, 3.3489],
[-1.0532, 3.4146],
[-0.9758, 3.6891],
[-1.0432, 3.267],
[-1.0349, 3.4871],
[-0.9621, 3.4335],
[-0.9611, 3.1801],
[-0.9632, 3.4894],
[-0.9695, 3.378],
[-1.0425, 3.4522],
[-0.951 , 3.6104],
[-0.8351, 3.1463],
[-0.7765, 3.3949],
[-1.0098, 3.7191],
[-1.1122, 3.7663],
[-1.0118, 3.5502],
[-0.8661, 3.7044],
[-0.9247, 3.7037],
[-0.985 , 3.7465]])

Matplotlib cannot handle categorical variables directly. Thus, we encode every class as a number so that we can incorporate the class labels into our plot.

```
1 le = LabelEncoder()  
2 y = le.fit_transform(df['class'])
```

```
1 plt.xlabel('LD1')  
2 plt.ylabel('LD2')  
3 plt.scatter(  
4     X_lda[:,0],  
5     X_lda[:,1],  
6     c=y,  
7     cmap='rainbow',  
8     alpha=0.7,  
9     edgecolors='b'  
10 )
```

<matplotlib.collections.PathCollection at 0x7f3d7323ad68>



Rather than implementing the Linear Discriminant Analysis algorithm from scratch every time, we can use the predefined LinearDiscriminantAnalysis class made available to us by the scikit-learn library.

```
1 from sklearn.discriminant_analysis import LinearDiscriminantAnalysis  
2 lda = LinearDiscriminantAnalysis()  
3 X_lda = lda.fit_transform(X, y)
```

```
1 lda.explained_variance_ratio_
```

array([0.6875, 0.3125])

```
1 plt.xlabel('LD1')  
2 plt.ylabel('LD2')  
3 plt.scatter(  
4     X_lda[:,0],  
5     X_lda[:,1],  
6     c=y,  
7     cmap='rainbow',  
8     alpha=0.7,  
9     edgecolors='b'  
10 )
```

