


Principal Component Analysis including Cumulative Explained Variance (Iris and MNIST datasets)

```
1  # Code source: Gaël Varoquaux
2  # Modified for documentation by Jaques Grobler
3  # License: BSD 3 clause
4
5  import matplotlib.pyplot as plt
6  from mpl_toolkits.mplot3d import Axes3D
7  from sklearn import datasets
8  from sklearn.decomposition import PCA
9
10 # import some data to play with
11 iris = datasets.load_iris()
12 X = iris.data[:, :2] # we only take the first two features.
13 y = iris.target
14
15 x_min, x_max = X[:, 0].min() - .5, X[:, 0].max() + .5
16 y_min, y_max = X[:, 1].min() - .5, X[:, 1].max() + .5
17
18 plt.figure(2, figsize=(8, 6))
19 plt.clf()
20
21 # Plot the training points
22 plt.scatter(X[:, 0], X[:, 1], c=y, cmap=plt.cm.Set1,
23             edgecolor='k')
24 plt.xlabel('Sepal length')
25 plt.ylabel('Sepal width')
26
27 plt.xlim(x_min, x_max)
28 plt.ylim(y_min, y_max)
29 plt.xticks(())
30 plt.yticks(())
31
32 # To get a better understanding of interaction of the dimensions
33 # plot the first three PCA dimensions
34 fig = plt.figure(1, figsize=(8, 6))
35 ax = Axes3D(fig, elev=-150, azimuth=110)
36 X_reduced = PCA(n_components=3).fit_transform(iris.data)
37 ax.scatter(X_reduced[:, 0], X_reduced[:, 1], X_reduced[:, 2], c=y,
38            cmap=plt.cm.Set1, edgecolor='k', s=40)
39 ax.set_title("First three PCA directions")
40 ax.set_xlabel("1st eigenvector")
41 ax.w_xaxis.set_ticklabels([])
42 ax.set_ylabel("2nd eigenvector")
43 ax.w_yaxis.set_ticklabels([])
44 ax.set_zlabel("3rd eigenvector")
45 ax.w_zaxis.set_ticklabels([])
46
47 plt.show()
```





```
1 from sklearn.datasets import fetch_openml
2 mnist = fetch_openml('mnist_784')
```

```
1 from sklearn.model_selection import train_test_split
2 # test_size: what proportion of original data is used for test set
3 train_img, test_img, train_lbl, test_lbl = train_test_split(mnist.data, mnist.target, test_size=1/7.0, random_state=42)
```

```
1 from sklearn.preprocessing import StandardScaler
2 scaler = StandardScaler()
3 # Fit on training set only.
4 scaler.fit(train_img)
5 # Apply transform to both the training set and the test set.
6 train_img = scaler.transform(train_img)
7 test_img = scaler.transform(test_img)
```

```
1 from sklearn.decomposition import PCA
2 # Make an instance of the Model
3 pca = PCA(.95)
```

```
1 pca.fit(train_img)
```

```
➔ PCA(copy=True, iterated_power='auto', n_components=0.95, random_state=None,
    svd_solver='auto', tol=0.0, whiten=False)
```

```
1 train_img = pca.transform(train_img)
2 test_img = pca.transform(test_img)
```

```
1 from sklearn.linear_model import LogisticRegression
```

```
1 # all parameters not specified are set to their defaults
2 # default solver is incredibly slow which is why it was changed to 'lbfgs'
3 logisticRegr = LogisticRegression(solver = 'lbfgs')
```

```
1 logisticRegr.fit(train_img, train_lbl)
```

```
➔ /usr/local/lib/python3.6/dist-packages/sklearn/linear_model/_logistic.py:940: ConvergenceWarning: lbfgs failed to converge. Increase the number of iterations.
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG)
LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
                    intercept_scaling=1, l1_ratio=None, max_iter=100,
                    multi_class='auto', n_jobs=None, penalty='l2',
                    random_state=None, solver='lbfgs', tol=0.0001, verbose=0,
                    warm_start=False)
```

```
1 # Predict for One Observation (image)
2 logisticRegr.predict(test_img[0].reshape(1,-1))
```

```
➔ array(['0'], dtype=object)
```

```
1 # Predict for One Observation (image)
2 logisticRegr.predict(test_img[0:10])
```

```
➔ array(['0', '4', '1', '2', '4', '7', '7', '1', '1', '7'], dtype=object)
```

```
1 print('Accuracy Score: %0.3f' % logisticRegr.score(test_img, test_lbl))
```

```
➔
```

```
1 print('Estimated number of components: \n', (pca.n_components_))
```

Estimated number of components:
327

```
1 print('Amount of variance explained by each of the selected components: \n', (pca.explained_variance_))
```

Amount of variance explained by each of the selected components:

40.59415525	29.01680964	26.87224996	20.86433884	18.05674875	15.74196929
13.77127266	12.54612494	11.00094033	10.03492632	9.64585917	8.64956409
7.9968328	7.8424034	7.3814671	7.16859604	6.68823549	6.60819348
6.40138235	6.2331149	5.91955064	5.73838634	5.4859824	5.32445389
5.14807896	4.9699168	4.92104344	4.75373005	4.51349763	4.41052839
4.31141958	4.23396877	4.0886635	4.07135857	4.04759614	3.98510077
3.83642443	3.80463967	3.70089599	3.65235589	3.47094549	3.41322949
3.39367785	3.27195748	3.21237257	3.20305992	3.15633079	3.14508318
3.08792167	3.07664551	2.98364544	2.88352378	2.86964419	2.8042245
2.78942547	2.766402	2.68205915	2.6503738	2.630463	2.59839835
2.53415884	2.49134621	2.4813325	2.45818633	2.4338657	2.40472015
2.37029837	2.30997244	2.28731369	2.25720798	2.2316704	2.20791417
2.19627765	2.17658648	2.1622007	2.13263816	2.11856585	2.09915978
2.07887268	2.06769352	2.04099226	2.03129133	2.01085963	2.00592863
2.00189849	1.99304477	1.97483524	1.9633447	1.93452927	1.91554913
1.9036724	1.88455667	1.87426626	1.87028656	1.85508397	1.81058262
1.80477181	1.79400846	1.76726368	1.75065918	1.72635723	1.70760674
1.69206454	1.65908045	1.64486947	1.61970796	1.60841903	1.59814954
1.58567607	1.55547553	1.5540458	1.54411919	1.53421886	1.51709899
1.4838185	1.47340942	1.46966777	1.45220574	1.44236705	1.41238479
1.40454052	1.39277132	1.3734908	1.37171373	1.36062348	1.35284157
1.34119435	1.32909986	1.31466917	1.29719283	1.28413076	1.27598694
1.25762271	1.25412644	1.22615998	1.2037214	1.20262191	1.19642729
1.18147374	1.17994621	1.16277097	1.15244952	1.13700231	1.12554587
1.1221124	1.11309523	1.10683259	1.10497361	1.08577893	1.07759413
1.06398967	1.0607537	1.05401173	1.04887655	1.03275367	1.0286867
1.02412814	1.02000472	1.00950214	1.00724437	1.00169961	0.99974496
0.99888706	0.99578642	0.98456272	0.98066593	0.97476626	0.96868066
0.9647658	0.96287354	0.95242638	0.93805526	0.9347652	0.9291984
0.92557676	0.91081254	0.90771945	0.89715684	0.89321422	0.87778269
0.87064796	0.86800266	0.85784502	0.85514628	0.84663193	0.84059328
0.83311669	0.82824188	0.82113901	0.80726691	0.80012303	0.7987441
0.79259406	0.7745348	0.76589166	0.757824	0.75413251	0.74949348
0.74610482	0.74080133	0.73195497	0.7210011	0.71744195	0.70543181
0.70182957	0.69836981	0.69548244	0.69272035	0.6812968	0.67814608
0.67127001	0.66761094	0.65178243	0.64876063	0.64320378	0.63619284
0.63160897	0.62538878	0.62459465	0.61452465	0.61254544	0.60556023
0.6003223	0.59550632	0.59229828	0.58819919	0.5784161	0.57667491
0.57094249	0.56848965	0.56353296	0.56040965	0.55743955	0.55105236
0.5446728	0.5425334	0.5314415	0.5301174	0.52234868	0.52065704
0.51543481	0.50432032	0.50061778	0.49998156	0.49465156	0.49199992
0.49122325	0.48934319	0.48471085	0.47606919	0.47046353	0.46928044
0.46416894	0.45827886	0.45619664	0.4541791	0.44896345	0.44707118
0.44345598	0.4393925	0.43652367	0.43274673	0.43109523	0.42784114
0.42133268	0.41800004	0.41717605	0.4120701	0.40959595	0.40696498
0.40040754	0.39755227	0.39582781	0.39314513	0.38960975	0.38654351
0.38612122	0.38208062	0.38021094	0.37561398	0.36962957	0.36800565
0.36585549	0.36334019	0.36149871	0.35950529	0.35771871	0.35363913
0.35061089	0.34742429	0.34539619	0.34416414	0.34164758	0.33387059
0.33185349	0.32878938	0.32327587	0.31870139	0.31658394	0.31555253
0.31172425	0.31129761	0.3101868	0.30903872	0.30359758	0.30122891
0.29892752	0.29772821	0.29447653	0.29377227	0.29051022	0.28788908
0.28345726	0.28217358	0.27989973	0.27796788	0.27526011	0.27349215
0.27298909	0.27008938	0.26892642	0.26689507	0.266189	0.26287416
0.26188522	0.25859412	0.25807053]			

```
1 print('Principal axis representing the directions of maximum variance: \n', (pca.components_))
```

Principal axis representing the directions of maximum variance:

[-1.49613708e-19	-2.77555756e-17	-0.00000000e+00	...	-0.00000000e+00	-0.00000000e+00
5.47306713e-19	2.77555756e-16	1.11022302e-16	...	-0.00000000e+00	-0.00000000e+00
-1.16585286e-19	-5.55111512e-17	-1.38777878e-17	...	-0.00000000e+00	-0.00000000e+00
1.15162705e-18	9.54097912e-17	6.59194921e-17	...	0.00000000e+00	0.00000000e+00
-1.78512164e-18	-6.07153217e-17	6.07153217e-17	...	0.00000000e+00	0.00000000e+00
-1.17426323e-19	-1.40295761e-16	-9.29161262e-17	...	0.00000000e+00	0.00000000e+00]

```
1 print('Percentage variance explained by each of the selected components: \n', pca.explained_variance_ratio_)
```

Percentage variance explained by each of the selected components:

[0.05685361	0.04063911	0.03763558	0.02922128	0.02528914	0.02204721
0.01928718	0.01757131	0.01540722	0.01405428	0.01350938	0.01211403
0.01119986	0.01098358	0.01033802	0.01003988	0.00936712	0.00925502
0.00896537	0.00872971	0.00829055	0.00803682	0.00768332	0.00745709
0.00721007	0.00696055	0.0068921	0.00665777	0.00632132	0.00617711
0.0060383	0.00592983	0.00572632	0.00570209	0.00566881	0.00558128
0.00537305	0.00532854	0.00518324	0.00511526	0.00486119	0.00478035
0.00475297	0.0045825	0.00449905	0.004486	0.00442056	0.0044048
0.00432475	0.00430896	0.00417871	0.00403848	0.00401904	0.00392742
0.00390669	0.00387445	0.00375632	0.00371195	0.00368406	0.00363915
0.00354918	0.00348922	0.0034752	0.00344278	0.00340872	0.0033679
0.00331969	0.0032352	0.00320347	0.0031613	0.00312554	0.00309227
0.00307597	0.00304839	0.00302824	0.00298684	0.00296713	0.00293995
0.00291154	0.00289588	0.00285848	0.0028449	0.00281628	0.00280938
0.00280373	0.00279133	0.00276583	0.00274974	0.00270938	0.0026828
0.00266616	0.00263939	0.00262498	0.00261941	0.00259811	0.00253579
0.00252765	0.00251258	0.00247512	0.00245186	0.00241783	0.00239157
0.0023698	0.0023236	0.0023037	0.00226846	0.00225265	0.00223827
0.0022208	0.0021785	0.0021765	0.0021626	0.00214873	0.00212475
0.00207814	0.00206356	0.00205832	0.00203387	0.00202009	0.0019781
0.00196711	0.00195063	0.00192362	0.00192114	0.0019056	0.0018947
0.00187839	0.00186145	0.00184124	0.00181677	0.00179847	0.00178707
0.00176135	0.00175645	0.00171728	0.00168586	0.00168432	0.00167564
0.0016547	0.00165256	0.0016285	0.00161405	0.00159241	0.00157637
0.00157156	0.00155893	0.00155016	0.00154756	0.00152067	0.00150921
0.00149016	0.00148562	0.00147618	0.00146899	0.00144641	0.00144071
0.00143433	0.00142855	0.00141384	0.00141068	0.00140292	0.00140018
0.00139898	0.00139464	0.00137892	0.00137346	0.0013652	0.00135667
0.00135119	0.00134854	0.00133391	0.00131378	0.00130917	0.00130138
0.0012963	0.00127563	0.00127129	0.0012565	0.00125098	0.00122937
0.00121937	0.00121567	0.00120144	0.00119766	0.00118574	0.00117728
0.00116681	0.00115998	0.00115004	0.00113061	0.0011206	0.00111867
0.00111006	0.00108476	0.00107266	0.00106136	0.00105619	0.00104969
0.00104495	0.00103752	0.00102513	0.00100979	0.0010048	0.00098798
0.00098294	0.00097809	0.00097405	0.00097018	0.00095418	0.00094977
0.00094014	0.00093501	0.00091285	0.00090861	0.00090083	0.00089101
0.00088459	0.00087588	0.00087477	0.00086066	0.00085789	0.00084811
0.00084077	0.00083403	0.00082954	0.00082379	0.00081009	0.00080765
0.00079963	0.00079619	0.00078925	0.00078487	0.00078071	0.00077177
0.00076283	0.00075984	0.0007443	0.00074245	0.00073157	0.0007292
0.00072189	0.00070632	0.00070113	0.00070024	0.00069278	0.00068906
0.00068798	0.00068534	0.00067886	0.00066675	0.0006589	0.00065724
0.00065009	0.00064184	0.00063892	0.00063609	0.00062879	0.00062614
0.00062108	0.00061539	0.00061137	0.00060608	0.00060376	0.00059921
0.00059009	0.00058542	0.00058427	0.00057712	0.00057365	0.00056997
0.00056079	0.00055679	0.00055437	0.00055061	0.00054566	0.00054137
0.00054078	0.00053512	0.0005325	0.00052606	0.00051768	0.00051541
0.00051239	0.00050887	0.00050629	0.0005035	0.000501	0.00049528
0.00049104	0.00048658	0.00048374	0.00048201	0.00047849	0.0004676
0.00046477	0.00046048	0.00045276	0.00044635	0.00044339	0.00044194
0.00043658	0.00043598	0.00043443	0.00043282	0.0004252	0.00042188
0.00041866	0.00041698	0.00041243	0.00041144	0.00040687	0.0004032

```
1 import numpy as np
2 import matplotlib.pyplot as plt
3 %matplotlib inline
4 plt.plot(np.cumsum(pca.explained_variance_ratio_))
5 plt.title('Cumulative Explained Variance')
6 plt.xlabel('number of components')
7 plt.ylabel('cumulative explained variance');
```

